Congratulations! You passed!

Grade received 90.47%
Latest Submission Grade 90.48%
To pass 80% or higher

Go to next item

You're viewing this assessment in its original language
You can switch back to view this content in your preferred translation if you'd prefer. You won't lose any progress if you change languages. Show Spanish version

1. Using Image Generator, how do you label images? 1 / 1 point

- (●) It's based on the directory the image is contained in
- ( ) It's based on the file name
- ( ) TensorFlow figures it out from the contents
- ( ) You have to manually do it

Correct
That's right! The directory of the image is the label.

2. What method on the Image Generator is used to normalize the image? 1 / 1 point

- ( ) normalize_image
- ( ) normalize
- ( ) Rescale_image
- (●) rescale

Correct
You've got it! This is the correct method for normalizing images.

3. How did we specify the training size for the images? 1 / 1 point

- ( ) The training_size parameter on the validation generator
- ( ) The target_size parameter on the validation generator
- (●) The target_size parameter on the training generator
- ( ) The training_size parameter on the training generator

Correct
Exactly! target_size specifies the image training size

4. When we specify the input_shape to be (300, 300, 3), what does that mean? 1 / 1 point

- (●) Every Image will be 300x300 pixels, with 3 bytes to define color
- ( ) Every Image will be 300x300 pixels, and there should be 3 Convolutional Layers
- ( ) There will be 300 horses and 300 humans, loaded in batches of 3
- ( ) There will be 300 images, each size 300, loaded in batches of 3

Correct
Nailed it! input_shape specifies image resolution.

5. If your training data is close to 1.000 accuracy, but your validation data isn't, what's the risk here? 1 / 1 point

- ( ) You're underfitting on your validation data
- (●) You're overfitting on your training data
- ( ) No risk, that's a great result
- ( ) You're overfitting on your validation data

Correct

Great job! The analysis corresponds too closely to the training data, and may therefore fail to fit additional data.

6. Convolutional Neural Networks are better for classifying images like horses and humans because: **0.3333333333333333 / 1 point**

- [ ] There's a wide variety of horses
- [ ] There's a wide variety of humans
- [x] In these images, the features may be in different parts of the frame

Correct
Correct! The receptive fields of different neurons partially overlap such that they cover the entire visual field.

You didn't select all the correct answers

7. After reducing the size of the images, the training results were different. Why? **1 / 1 point**

- ( ) There was more condensed information in the images
- (•) We removed some convolutions to handle the smaller images
- ( ) There was less information in the images
- ( ) The training was faster

Correct
Yes! Removing some convolutions modifies the training results.