

Your grade: 100%

Next item →

Your latest: **100%** • Your highest: **100%** • To pass you need at least 80%. We keep your highest score.

i You're viewing this assessment in its original language
You can switch back to view this content in your preferred translation if you'd prefer. You won't lose any progress if you change languages. [Show Spanish version](#) ×

1. To use word embeddings in TensorFlow, in a Sequential model, what is the name of the layer? **1 / 1 point**

- ☐ tf.keras.layers.Word2Vector
- ☐ tf.keras.layers.Embed
- ☒ tf.keras.layers.Embedding
- ☐ tf.keras.layers.WordEmbedding

✓ **Correct**
That's right!

2. Using the default settings, what does the 'max_tokens' parameter do when initializing the TextVectorization layer? **1 / 1 point**

- ☒ It specifies the maximum size of the vocabulary, and picks the most common 'max_tokens - 2' words
- ☐ It specifies the maximum size of the vocabulary, and picks the most common 'max_tokens' words
- ☐ It specifies the maximum size of the vocabulary, and picks the most common 'max_tokens - 1' words
- ☐ It errors out if there are more than max_tokens distinct words in the corpus

✓ **Correct**
That's right! This includes the 0 index for padding and 1 for out-of-vocabulary words.

3. What is the name of the TensorFlow library containing common data that you can use to train and test neural networks? **1 / 1 point**

- ☐ TensorFlow Data Libraries
- ☐ TensorFlow Data
- ☒ TensorFlow Datasets
- ☐ There is no library of common data sets, you have to use your own

✓ **Correct**
Correct!

4. What is the purpose of the embedding dimension? **1 / 1 point**

- ☐ It is the number of letters in the word, denoting the size of the encoding
- ☐ It is the number of words to encode in the embedding
- ☒ It is the number of dimensions for the vector representing the word encoding
- ☐ It is the number of dimensions required to encode every word in the corpus

✓ **Correct**
That's right!

5. IMDB Reviews are either positive or negative. What type of loss function should be used in this scenario? **1 / 1 point**

- ☐ Adam
- ☒ Binary crossentropy
- ☐ Categorical crossentropy
- ☐ Binary Gradient descent

✓ **Correct**
Correct!

6. How are the labels for the IMDB dataset encoded?

1 / 1 point

- ☐ Reviews encoded as a boolean true/false
- ☒ Reviews encoded as a number 0-1
- ☐ Reviews encoded as a number 1-10
- ☐ Reviews encoded as a number 1-5

✓ Correct
Correct!

7. How many reviews are there in the IMDB dataset and how are they split?

1 / 1 point

- ☒ 50,000 records, 50/50 train/test split
- ☐ 60,000 records, 80/20 train/test split
- ☐ 60,000 records, 50/50 train/test split
- ☐ 50,000 records, 80/20 train/test split

✓ Correct
That's right!